The Verification Gap: Checking the Reasoning in AI-Assisted Academic Writing

Jeroen Veen[1]

[1]HAN University of Applied Sciences, Arnhem and Nijmegen, the Netherlands

jeroen.veen@han.nl

ORCID: 0009-0008-5892-5782

paper1-draft-4

Key Points:

- Generative AI tools can produce fluent papers whose citations are real and whose reasoning is still weak.
- Reporting guidelines help authors report a study completely; they do not ask whether each step of reasoning holds.
- A claim register records what each claim rests on and checks each step by its kind, with wording tied to support.
- Each part of the register is used elsewhere; none of the approaches compared here combines them for academic papers.
- Publishing a register can make reasoning traceable; whether it helps reviewers find weak steps is untested.

1 Introduction

Writing assistants based on large language models (LLMs) have spread quickly: up to 22% of computer science papers now show signs of LLM editing, and such editing can be detected in every major discipline (Liang et al., 2025). In this paper "generative AI tools", or "AI" for short, means these assistants. When a human author falls short, the text can show it, in awkward phrasing or gaps in coverage. An AI-assisted draft can read as polished and complete while it contains errors that only checking would reveal.

The most discussed of these errors are citation errors: a reference to a work that does not exist, or to a real work with the wrong details. Mugaanyi et al. (2024) found that 62–89% of the digital object identifiers (DOIs) that ChatGPT (GPT-3.5) generated for natural-science and humanities topics were fabricated. Walters and Wilder (2023) found fabricated citations among the references of generated literature reviews, with fewer from the newer of the two models they tested. A 2026 benchmark tested newer models that can search the web. Asked for the bibliographic entries of 931 named papers, three of them (GPT-5, Claude Sonnet 4.6 and Gemini 3 Flash) got only 50.9% of the entries fully right; most errors were in single fields, but some were details of a different paper (Rao and Callison-Burch, 2026, §4).

A second error cannot be caught by a lookup. I call it *confidence inflation*: wording that expresses more certainty than the evidence supports. The evidence indicates that LLMs tested in 2023 and 2024 verbalised high confidence even on questions they answered wrongly (Xiong et al., 2024; Zhou et al., 2024). Whether this carries over into academic prose is unsettled: in one corpus of applied-linguistics abstracts, ChatGPT used boosters such as "shows" at twice the human rate, but also more hedges (El-Dakhs et al., 2025). The stakes are epistemic, because words such as "demonstrates" and "suggests" are meant to tell a reader how much weight a claim can bear.

A third error is harder to see, because each citation passes: the sources are real and say what is attributed to them, but the claim drawn from them says more. This paper's own register first stated that LLMs "invent plausible-sounding papers", in the present tense, from two studies of ChatGPT models from 2023 and 2024. Each study was cited

correctly; the step from them to all current LLMs was not supported. A reader caught it, not a citation check. Checking the register's source-based claims against their sources' scope then found three more like it (Veen, 2026).

A citation check, as I use the term and as such checks are usually run, asks whether a cited work exists and says what is attributed to it, so it catches the first error. It does not reach the second or the third (Section 4 addresses both): a claim can rest on correct citations and still say more than they carry, so a fluent paper can pass the check with its reasoning unchecked.

Detecting AI-written text does not fill that gap. Detection asks who wrote a text; verification asks whether it holds. Detectors may now answer the first question well for pure cases: one 2025 working paper found that the best of three commercial detectors separated published human prose from fully generated text almost without error (Jabarian and Imas, 2025). Text that a person directs and an AI drafts, as this paper was, was not tested; a detector still gives a verdict on such text, but the verdict concerns the drafting, not whether the text holds. But journal guidance now accepts AI help that authors disclose: the ICMJE recommendations ask authors to declare AI use and describe how they used it (International Committee of Medical Journal Editors, 2026). Once AI help is allowed and declared, the question that remains is whether the text holds.

2 Reporting Guidelines

The most developed infrastructure for checking academic writing is the EQUATOR Network's collection of about 700 reporting guidelines for health research (EQUATOR Network, n.d.). Guidelines such as CONSORT for randomised trials and PRISMA for systematic reviews (Page et al., 2021) give authors a checklist to complete while they prepare a manuscript. The GoodReports tool picks the checklist that fits a study type (Struthers et al., 2021).

What these guidelines check is reporting. Their developers define a reporting guideline as "a checklist, flow diagram, or explicit text to guide authors in reporting a specific type of research" (Moher et al., 2010). For CONSORT, the measured effect is on completeness:

journals endorsing CONSORT reported more completely on 25 of 27 checklist items, significantly so for five (Turner et al., 2012). Incomplete reporting is a recognised source of research waste (Glasziou et al., 2014), so this is a real service. However, the question it answers is not whether the reasoning holds. That question applies to any paper, and matters most where a paper has little else to check.

Gregor (2006) distinguishes five types of theory in information systems research, among them analytic theories (Type I) and theories for design and action (Type V). When a paper of either type proposes a new theory, no earlier source states it, because the proposal is the contribution. I call papers whose main contribution is reasoning *non-empirical*: theoretical, conceptual, design-science and perspective papers are examples.

In a non-empirical paper the reasoning carries the contribution. Toulmin (2003) places the step from evidence to claim in the warrant, and Whetten (1989, p. 491) calls the reasoning of a theory "the theoretical glue that welds the model together". A check that stops at citations may therefore leave the paper's contribution unchecked.

Guidance for papers without new data exists, but it is scattered. Conceptual papers lack "commonly accepted templates" (Jaakkola, 2020, p. 18), although Jaakkola proposes four. Design-science researchers have reported no agreed guidance on presenting their work (Gregor and Hevner, 2013; Dinter and Krawatzeck, 2015), and validity in that field "is neither well-understood nor consistently established" (Larsen et al., 2025). Software engineering's community standards ask methodological papers for "clear, valid arguments" but give no procedure for checking them (Ralph et al., 2020).

To look further, I searched for a consensus-based checklist that checks the reasoning of non-empirical papers step by step: whether each step from evidence to claim holds. On 16 October 2026 I searched the EQUATOR library (709 titles) and four bibliographic sources (Europe PMC, OpenAlex, Crossref and arXiv) with 28 queries, and I screened 701 titles alone. I found none. The search did not cover Scopus, Web of Science or humanities indexes, and a second screener has not yet checked it. The protocol, queries and screened candidates are public (Veen, 2026), so the search can be repeated.

3 Related Approaches

The register proposed in Section 4 has four parts. Each claim is recorded with its grounds, and each kind of grounds gets its own check. The step from grounds to claim is checked for relevance and sufficiency. The wording of a claim is tied to the strength of its support. Several practices already have one or more of these parts (Table 1).

Table 1: Which parts of the claim register (Section 4) each approach has: *yes*, *partly* (see the notes) or –. The cells are this paper's reading of each source.

Approach	**Scope**	**Claims with grounds**	**Check by kind of grounds**	**Step checked**	**Wording tied to support**
Reporting guidelines (Moher et al., 2010)	One type of empirical study	–	–	–	–
AI-writing checkers (Samsonau, 2026; Chen et al., 2026; Zhou and Yu, 2026)	AI-assisted manuscripts	yes[a]	–	partly[b]	partly[c]
Appraisal tools (JBI, 2026; Baethge et al., 2019)	Expert opinion (JBI); narrative reviews (SANRA)	–	–	partly[d]	–
Assurance cases (Kelly and Weaver, 2004; Bloomfield and Netkachova, 2014)	Safety of a system	yes	–	partly[e]	–
Micro- and nanopublications (Clark et al., 2014; Groth et al., 2010)	Statements in (mainly biomedical) publications	yes	–	–	–
Argument-based validation (Kane, 2013)	Interpretation and use of test scores	yes	partly[f]	yes	–
Argument quality (Wachsmuth et al., 2017)	Arguments in natural language	–	–	partly[g]	–
Calibrated language (Kent, 1964; Mastrandrea et al., 2010; Santesso et al., 2020)	Estimates, assessments, evidence summaries	–	–	–	yes
Claim register (Section 4)	Academic papers	yes	yes	yes	yes

[a] Each claim is paired with an evidence packet (Chen et al., 2026). [b] Whether the evidence supports the claim (Chen et al., 2026) or, on a sample, suffices (Zhou and Yu, 2026); not a check of each claim's step. [c] Causal wording checked against the study design (Samsonau, 2026). [d] One question or item on the argument as a whole. [e] Arguments are stated, each with a side-warrant (Bloomfield and Netkachova, 2014); whether the case convinces is left to judgement. [f] "Different kinds of warrants require different kinds of backing" (Kane, 2013, p. 14): by kind of warrant, not of grounds. [g] Criteria (acceptability, relevance, sufficiency), not a procedure for papers.

Recent tools for AI-assisted writing go beyond checking that references exist. One linter also checks causal language against the study design (Samsonau, 2026). One workflow returns unsupported claims to the author (Chen et al., 2026), and a provenance discipline samples whether the evidence suffices (Zhou and Yu, 2026). A checklist for cited assertions in narrative reviews is in development (VANRA Development Group, 2026). As their authors describe them, each asks whether evidence exists and supports a claim; Zhou and Yu judge sufficiency on a sample, and their tiers rate metrics. None describes a check, for each claim in the prose, of whether its grounds bear on it and justify its wording, and what would make the step fail.

The appraisal tools come closer to the reasoning. The evidence-synthesis organisation JBI has a consensus-based checklist for appraising expert opinion; it asks "Does the opinion demonstrate a logically defended argument to support the conclusions drawn?" and refers reviewers to Toulmin's model (McArthur et al., 2025; JBI, 2026). SANRA, a scale for narrative reviews, scores scientific reasoning as one of its six items, and its authors call it "not a reporting guideline" (Baethge et al., 2019). Reporting items for concept analysis have only recently been collected (Nashwan et al., 2026). Both judge an argument as a whole.

Two practices record the structure of an argument explicitly. In safety engineering, an assurance case, or safety case, is "a documented body of evidence that provides a convincing and valid argument that a system is adequately safe" for its use (Bishop and Bloomfield, 1998). Its authors "make an explicit set of claims", link them to the evidence through arguments, and "make clear the assumptions and judgements underlying the arguments" (Bishop and Bloomfield, 1998). The Goal Structuring Notation breaks claims into sub-claims until each can be supported by evidence (Kelly and Weaver, 2004), and the claims–arguments–evidence approach defines its arguments as "Toulmin's warrants" and treats claims asserted without justification as assumptions (Bloomfield and Netkachova, 2014).

In scholarly data publishing, nanopublications express a core scientific statement with its context and provenance (Groth et al., 2010). Micropublications record a statement

with its attribution and, at their fullest, with the evidence and challenges brought for and against it; their model "is grounded in Toulmin-Verheij" (Clark et al., 2014). Both make claims and their support machine-readable. Neither describes a procedure for checking a step.

Educational measurement has an example of checking an argument step by step. In Kane's argument-based approach, the claims made from test scores are laid out as an argument that specifies the inferences and assumptions they need. Validation then evaluates the coherence and completeness of that argument and the plausibility of its inferences and assumptions (Kane, 2013). The object is the use of a test, not a paper.

Research on argument quality supplies criteria for the step itself. Wachsmuth et al. (2017) distinguish logical, rhetorical and dialectical quality, and define a cogent argument as one with "individually acceptable premises that are relevant to the argument's conclusion and, together, sufficient to draw the conclusion". Walton et al. (2008) collect argumentation schemes, which their publisher describes as "forms of argument that capture stereotypical patterns of human reasoning".

Each of the four parts therefore already has users (Table 1), but none of the approaches compared here combines them for an academic paper. That combination is what the next section proposes.

4 The Claim Register

This paper proposes a claim register: a table, published with the paper, with one row for each claim and one for the paper's contribution. Toulmin's layout of arguments gives each claim's row and its step-by-step check; Whetten's account of a theoretical contribution gives the contribution's row, checked once as a whole. Box 1 gives the terms; the framework's glossary defines them from their sources (Veen, 2026).

Toulmin modelled his layout of arguments on jurisprudence, calling logic "generalised jurisprudence" (Toulmin, 2003, p. 7). In his account any assertion is a claim. Its *data* are "the facts we appeal to as a foundation for the claim". Its warrant is one of the "general, hypothetical statements, which can act as bridges" from data to claim. A

Box 1. Terms used in this paper.

Term	Meaning
Claim register	The table, published with the paper, that records each claim with its grounds, check and tier, and the contribution with its check.
Claim	Any assertion put forward for acceptance.
Grounds	What a claim rests on: a source, other claims, a declared starting point, or the author's own work. An extension of Toulmin's *data*.
Warrant	The step from grounds to claim.
Qualifier	The hedge that states how strong the claim is.
Rebuttal	Circumstances in which the step from grounds to claim no longer holds.
Confidence tier	The rating (Established, Supported, Emerging, Speculative) that decides which qualifiers a claim may carry.
Priority	How much the paper depends on a claim (P0 to P2).
Declared starting point	An assumption, definition or value the paper adopts without arguing for it.
Contribution	What the paper adds; checked once, as a whole.
Verification	Checking each claim against its grounds and warrant.
Validation	Checking that the contribution serves its readers and purpose.

qualifier indicates "the strength conferred by the warrant on this step"; and a *rebuttal* names "circumstances in which the general authority of the warrant would have to be set aside" (Toulmin, 2003, pp. 90–94). Toulmin later wrote that he had not had in mind an analytical model like the one that came to be called "the Toulmin model"; his concern "was with twentieth-century epistemology, not informal logic" (Toulmin, 2003, p. vii). Using his layout as a checklist is my adaptation, and it has to be judged on its own merits.

Others have put his elements to work. Micropublications build on his layout to record the claims of (mainly biomedical) publications (Clark et al., 2014). Kane's argument-based validation analyses the inferences from test scores to their uses with his model (Kane, 2013, pp. 11–12). Two studies of university students' essays scored argument quality element by element, with two raters each (Yang, 2022; Divya and Alamelu, 2026). None of these uses checks a research paper's reasoning, so that application is a further step.

Checks by kind of grounds. In the register, each claim is recorded with its grounds item by item, and each kind of grounds gets its own check:

- a **source** is checked in two steps: the work exists, and it says what the claim attributes to it, for the cases and period the claim covers;

- a **declared starting point** is checked for being stated, with a note on whether the field accepts or contests it;
- the author's **own work**, such as a search or a trial, is checked for a method someone else could repeat;
- a step from **other claims**, or from several sources to a point none of them makes alone, gets the warrant check.

The warrant check. The warrant check asks two questions, taken from Wachsmuth and colleagues' definition of a cogent argument. *Relevance*: does each ground bear on this claim, so that the claim would be less credible without it? *Sufficiency*: do the grounds together justify the claim at the strength its qualifier states? A sufficiency check names what is missing, such as a case not covered, a search not run or a premise that rests on one study. The check also records the rebuttal, the circumstances in which the step would fail, and the strongest known objection with the answer to it. Where a step follows a common pattern, such as an argument from expert opinion or from the absence of search results, the questions proper to that pattern could be added. I have not yet worked them out from the argumentation-schemes literature.

The check has four outcomes: it passes, it passes and someone other than the author confirms it, it leaves a named gap open, or it fails. A failed check means that the recorded grounds do not carry the claim as worded, not that the claim is false; rewording the claim to what the grounds carry is the usual response.

Tiers and wording. Fixed wording for levels of certainty has precedents. Kent (1964) argued for "a few key odds expressions" in intelligence estimates. The Intergovernmental Panel on Climate Change (IPCC) asks its authors to judge confidence from the "type, amount, quality, and consistency of evidence" and "the degree of agreement" (Mastrandrea et al., 2010). The GRADE working group (Grading of Recommendations Assessment, Development and Evaluation) fixes wording per level of certainty and size of effect, from "results in" to "may result in" (Santesso et al., 2020). Fixed terms may not secure a

shared reading: when readers assigned numbers to the probability terms in an IPCC report, their judgements deviated from the IPCC's guidelines even when the guidelines were at hand (Budescu et al., 2009).

Table 2 adapts the practice for academic papers. A claim's tier is the lower of two ratings: how well its grounds are supported, and what its warrant check allows. An argued claim can therefore reach the Supported tier when its premises are supported and its check passes, which a rule that only counts sources does not allow. The table has not been tested with readers; it may help an author notice when a sentence claims more than its grounds allow.

Table 2: Confidence tiers and the qualifiers they permit (proposed). A claim takes the lower of the two ratings.

Tier	**Support for the grounds**	**Warrant check**	**Qualifiers**
Established	Three or more independent sources agree, at least two of them peer-reviewed; no known dispute	Passed and confirmed by someone other than the author	"shows", "confirms"
Supported	Two independent peer-reviewed sources agree	Passed	"indicates", "evidence suggests"
Emerging	One peer-reviewed source; two or more that are not; or pilot data	A named gap left open	"may", "preliminary evidence"
Speculative	No source reports the point, or a single one that is not peer-reviewed	Failed	"warrants investigation"

A claim about what one source says may reach Established (this paper holds such claims at Supported); searches and design trials are rated by method.

Priority. Every claim also carries a *priority*, the author's answer to the question of what fails if the claim is wrong. A P0 claim is one the paper fails without; P1 strengthens the argument; P2 gives context. Priority sets how much checking a claim deserves, and every P0 entry, the contribution included, should reach the Supported tier. A paper whose central claims stay at Emerging falls short of that floor; the register makes this visible.

The contribution. Besides its claims, the register has one entry for what the paper contributes, and that entry is checked differently: once, as a whole, rather than step by step. The check comes from Whetten. He set out what "a complete theory must

contain": which factors (What), how they relate (How), why ("the theoretical glue that welds the model together"), and for whom, where and when it holds (Whetten, 1989, pp. 490–492). His seven reviewer questions, from "What's new?" to "Who cares?", answer "What constitutes a publishable theory paper?" (Whetten, 1989, p. 494). They judge a contribution as a whole, so they may serve as a check of a paper's central contribution, applied once. For a contribution that is a recommendation, they fit less well: How tends to restate the thesis, and the check needs a declared value and the conditions under which the recommendation applies.

These two levels map onto a distinction from systems engineering. Verification is "proof of compliance with requirements"; validation shows that a product "accomplishes the intended purpose in the intended environment" (NASA, 2016, §2.0). Checking each claim against its grounds and warrant is verification in this sense; asking Whetten's questions of the contribution is closer to validation. The mapping is mine. The register then works as a traceability record, of the kind systems engineers keep for requirements. Each link is written down, from the contribution to its claims, from a claim to its grounds and from grounds to their sources, so that someone else can follow and check it. The IPCC asks its authors for a "traceable account" of their evaluation of evidence for the same reason (Mastrandrea et al., 2010).

What passing means. A register that passes its checks means that the paper followed the method: each claim's grounds and warrant were recorded and checked against stated criteria. It does not mean that the claims are true. A check run by the author, or by the author's AI tools, can err in the author's favour, which is why only a check confirmed by someone else can reach the Established tier.

The central argument of the proposal rests on these pieces. A citation check, or a rating that counts sources, does not look at the step from grounds to claim: it can pass a claim that says more than its sources carry, and fail a sound inference with no source to count. A check chosen by the kind of grounds (a source against the claim's scope, an inference against its qualifier) looks at that step and may avoid both errors. In this paper's register a scope check found three overreaches after a reader found the first (Section 1);

the remedy is untried on another author's paper.

5 Applying the Register to This Paper

This paper was written with the procedure it proposes. Its claim register, published with the paper (Veen, 2026), records each claim's grounds, warrant, warrant check, qualifier, confidence tier and priority, plus the paper's declared starting points and its one contribution. Table 3 shows three entries: a claim resting on sources, a claim resting on the author's own search, and the contribution.

Table 3: Three register entries: what each rests on, how it was checked, and how its priority and confidence tier were set.

	S1-1	**S2-2**	**C1 (contribution)**
Statement	AI citation hallucination is a distinct failure mode.	No consensus checklist was found for checking the reasoning of non-empirical papers step by step.	A published register makes reasoning traceable; registers should be trialled before they are adopted.
Says	descriptive	descriptive	normative
Grounds	two sources: Mugaanyi et al. (2024), Walters and Wilder (2023)	the author's documented search; gap statements for conceptual and design-science papers	four claims, two declared values, three declared definitions
Check	source check, both steps, for each source	method check of the search; warrant check, because absence is inferred from a search	Whetten's elements (Table 4); boundary conditions and values stated
Priority	P0: the opening argument fails without it	P0: the gap the paper addresses	P0 by definition: it is what the paper contributes
Tier	Supported	Emerging	Emerging
How it was set	two independent peer-reviewed sources agree	one screener and no search of Scopus, Web of Science or humanities indexes; the warrant check names the same gap	no higher than its weakest supporting claim (S2-2 and S3-4 are Emerging)
Qualifier in the text	reported findings ("found")	"I found none"	"I propose", "could", "should test"

What the register shows about this paper. Of the six P0 entries, three reach the Supported tier and three stay at Emerging: the gap claim (S2-2), the central argument (S3-4) and the contribution (C1), which rests on both. Under the rule in Table 2 they stay at Emerging for reasons that are named and could be removed: the search had one

Table 4: Whetten's elements applied to the contribution (C1).

Element	Content
What	A claim register; non-empirical papers; AI-assisted writing (each a declared definition).
How	Publishing the register exposes each step from grounds to claim to readers and reviewers.
Why	Citation checks pass weak reasoning (S1-7), reporting guidelines check reporting (S2-5), no checklist was found for the reasoning (S2-2), and each claim needs the check its grounds call for (S3-4).
Who, where, when	Papers whose contribution is an argument, and the interpretive parts of empirical papers; other writing, such as journalism, untested; frameworks from Western traditions.
What's new	The combination: each part is used elsewhere (S4-6 to S4-8, S3-6), and none of the compared approaches combines them for academic papers (S4-9).
Values	Published reasoning should be checkable (D1); a proposal should be tested before it is adopted (D7).
Objections	Peer review already judges arguments: a register would support review, not replace it. The register burdens authors: a shorter reviewer form is named as an option (Section 6). An AI can draft a register fluently: the register's author must check each entry by hand.

screener (S2-2), and the remedy was tried only by one author (S3-4). A second screener and a register kept by someone else are the evidence this paper still lacks.

6 Limitations

Its parts have users elsewhere (Section 3), but the register itself has one so far: the author of this paper. Whether it works in other hands is untested. Its reviews came from five language models from four vendors, in sessions separate from the drafting. They are independent of it but not of the method, and no substitute for human reviewers.

The register asks work of authors; this paper's has 36 entries. A shorter form, the warrant check's questions for a reviewer to ask of the central claims, may suit many papers better. Toulmin's and Whetten's frameworks come from Western academic traditions, and scholars in the humanities may find a claim-by-claim layout foreign to how their arguments work; their fit with other conventions needs to be assessed.

The checks are also partly circular: the author who wrote the rules applies them, and

the warrant check is a judgement, recorded with its reasons so that someone else can disagree with it. A register is itself text that an AI tool can draft fluently, as it drafted this paper's, so its author must check each entry by hand before others rely on it.

7 Recommendations

The value behind this proposal is that published reasoning should be checkable: readers and reviewers should be able to see what each claim rests on. A claim register, published with a paper whose contribution is an argument, makes that reasoning visible and traceable. I propose two steps, in this order. Authors and journals could publish registers as optional supplementary material, as this paper does, and trials should test whether reviewers who have a register find different weaknesses from reviewers who do not. A register should become part of verification infrastructure for non-empirical papers, of the kind EQUATOR offers for empirical health research, only if such trials find that it helps. A first implementation exists: templates for the register and the glossary, the checks for each kind of grounds, and tools that test a register against its manuscript (see Code and Data Availability).

The proposal applies to papers whose contribution is an argument, and to the interpretive parts of empirical papers. Other writing that asks readers to trust its reasoning, journalism among them, might use it too, but that is untested. Peer review already judges arguments, and a register would not replace it. It could show reviewers where a paper's weight rests, and it could help authors find weak steps before submission; whether it does either is the question the trials should answer.

Several further directions warrant investigation. The gap search needs a second screener and wider databases. The questions proper to common patterns of argument need working out, and the checks need adapting to other fields. The checks could also be built into writing tools, so that they run while an author drafts.

Acknowledgements and use of AI

The author conceived the idea and developed the framework, starting in March 2026, and made the decisions on scope, structure and the framework's rules. Generative AI was used throughout, under the author's direction. Claude Opus 5.5 (Anthropic), through Claude Code, drafted and revised the prose, searched for and checked sources, drafted the claim register and ran the framework's checking tools. The author directed the revisions and reviewed each draft. Five models from four vendors reviewed the first two drafts and the framework's glossary in sessions separate from the drafting: DeepSeek-V4-Pro, Gemini 3.1 Pro (preview), GPT-6 (OpenAI, through ChatGPT), and Claude Opus 5.5 and Claude Fable 5.1 (Anthropic). Their reviews informed this version. The author takes full responsibility for the content. This work received no specific funding. Verification record, run on this version on 30 October 2026: all 22 DOIs in the claim register resolve; of the 43 entries with a DOI in the bibliography file, 34 match their Crossref or DataCite record and 9 differ in minor fields (such as an author list cut short with "others"), with none failing; the register's consistency check finds no schema, anchor or tier errors, and eight claims rest on entries that await an independent re-rating; three of the six P0 entries reach the Supported tier (Section 5).

Conflict of interest

The author develops the open-source framework described here and receives no income from it; there are no financial conflicts.

Ethics statement

Not applicable: the work involved no human participants or personal data.

Code and data availability

The framework described here (templates, glossary and checking tools) and this paper's claim register are public under a CC BY 4.0 licence at `https://github.com/ducroq/agent-ready-papers` (Veen, 2026); the register is `papers/perspective/vv/claims/claim_registry.md`. The version this paper describes is tagged `paper1-preprint-1` in that repository.

References

Baethge, C., Goldbeck-Wood, S., and Mertens, S. (2019). SANRA—a scale for the quality assessment of narrative review articles. *Research Integrity and Peer Review*, 4:5.

Bishop, P. and Bloomfield, R. (1998). A methodology for safety case development. In Redmill, F. and Anderson, T., editors, *Industrial Perspectives of Safety-critical Systems*, pages 194–203. Springer London.

Bloomfield, R. and Netkachova, K. (2014). Building blocks for assurance cases. In *2014 IEEE International Symposium on Software Reliability Engineering Workshops*, pages 186–191.

Budescu, D. V., Broomell, S., and Por, H.-H. (2009). Improving communication of uncertainty in the reports of the intergovernmental panel on climate change. *Psychological Science*, 20(3):299–308.

Chen, G., Yu, Y., and Wang, W. (2026). Evidence-ledger adjudication for claim-evidence traceability.

Clark, T., Ciccarese, P. N., and Goble, C. A. (2014). Micropublications: A semantic model for claims, evidence, arguments and annotations in biomedical communications. *Journal of Biomedical Semantics*, 5:28.

Dinter, B. and Krawatzeck, R. (2015). Towards a configurative publication schema for

design science research. In *Proceedings of the 36th International Conference on Information Systems (ICIS 2015)*. Research-in-Progress.

Divya, T. J. and Alamelu, C. (2026). Integrating DeepL write and Claude AI to enhance argumentative writing competence of engineering students in India. *Frontiers in Artificial Intelligence*, 9:1838463.

El-Dakhs, D. A. S., Afzaal, M., and Siyanova-Chanturia, A. (2025). A genre-based comparison of Chat-GPT-generated abstracts versus human-authored abstracts: Focus on applied linguistics research articles. *Corpus Pragmatics*, 10(1):1.

EQUATOR Network (n.d.). Enhancing the QUAlity and Transparency of health research. `https://www.equator-network.org/`. Accessed: 2026-03-03.

Glasziou, P., Altman, D. G., Bossuyt, P., Boutron, I., Clarke, M., Julious, S., Michie, S., Moher, D., and Wager, E. (2014). Reducing waste from incomplete or unusable reports of biomedical research. *The Lancet*, 383(9913):267–276.

Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3):611–642.

Gregor, S. and Hevner, A. R. (2013). Positioning and presenting design science research for maximum impact. *MIS Quarterly*, 37(2):337–355.

Groth, P., Gibson, A., and Velterop, J. (2010). The anatomy of a nanopublication. *Information Services & Use*, 30(1–2):51–56.

International Committee of Medical Journal Editors (2026). Recommendations for the conduct, reporting, editing, and publication of scholarly work in medical journals, §V: Use of artificial intelligence in publishing. `https://www.icmje.org/recommendations/browse/artificial-intelligence/ai-use-by-authors.html`. Updated January 2026. Accessed: 2026-09-27.

Jaakkola, E. (2020). Designing conceptual articles: four approaches. *AMS Review*, 10(1–2):18–26.

Jabarian, B. and Imas, A. (2025). Artificial writing and automated detection. NBER Working Paper 34223, National Bureau of Economic Research.

JBI (2026). Checklist for textual evidence: Expert opinion. JBI critical appraisal tools.

Kane, M. T. (2013). Validating the interpretations and uses of test scores. *Journal of Educational Measurement*, 50(1):1–73.

Kelly, T. and Weaver, R. (2004). The goal structuring notation – a safety argument notation. In *Proceedings of the Dependable Systems and Networks 2004 Workshop on Assurance Cases*.

Kent, S. (1964). Words of estimative probability. *Studies in Intelligence*, 8(4).

Larsen, K. R., Lukyanenko, R., Mueller, R. M., Storey, V. C., Parsons, J., VanderMeer, D., and Hovorka, D. S. (2025). Validity in design science. *MIS Quarterly*, 49(4):1267–1294.

Liang, W., Zhang, Y., Wu, Z., Lepp, H., Ji, W., Zhao, X., Cao, H., Liu, S., He, S., Huang, Z., Yang, D., Potts, C., Manning, C. D., and Zou, J. Y. (2025). Quantifying large language model usage in scientific papers. *Nature Human Behaviour*, 9(12):2599–2609.

Mastrandrea, M. D., Field, C. B., Stocker, T. F., Edenhofer, O., Ebi, K. L., Frame, D. J., Held, H., Kriegler, E., Mach, K. J., Matschoss, P. R., Plattner, G.-K., Yohe, G. W., and Zwiers, F. W. (2010). Guidance note for lead authors of the IPCC fifth assessment report on consistent treatment of uncertainties. Technical report, Intergovernmental Panel on Climate Change (IPCC).

McArthur, A., Cooper, A., Edwards, D., Klugarova, J., Yan, H., Barber, B. V., Gregg, E. E., Weeks, L. E., and Jordan, Z. (2025). Textual evidence systematic reviews series paper 3: critical appraisal of evidence from narrative, opinion, and policy. *JBI Evidence Synthesis*, 23(5):833–839.

Moher, D., Schulz, K. F., Simera, I., and Altman, D. G. (2010). Guidance for developers of health research reporting guidelines. *PLoS Medicine*, 7(2):e1000217.

Mugaanyi, J., Cai, L., Cheng, S., Lu, C., and Huang, J. (2024). Evaluation of large language model performance and reliability for citations and references in scholarly writing: Cross-disciplinary study. *Journal of Medical Internet Research*, 26:e52935.

NASA (2016). *NASA Systems Engineering Handbook*. National Aeronautics and Space Administration.

Nashwan, A. J., Kunjavara, J., Joy, G., Singh, K., Rony, M. K. K., Mannethodi, K., and Ahmed, S. K. (2026). Preferred reporting items for concept analysis in nursing: a systematic review. *Journal of Research in Nursing*, 31(5–6):498–511.

Page, M. J., McKenzie, J. E., Bossuyt, P. M., Boutron, I., Hoffmann, T. C., Mulrow, C. D., Shamseer, L., Tetzlaff, J. M., Akl, E. A., Brennan, S. E., Chou, R., Glanville, J., Grimshaw, J. M., Hróbjartsson, A., Lalu, M. M., Li, T., Loder, E. W., Mayo-Wilson, E., McDonald, S., McGuinness, L. A., Stewart, L. A., Thomas, J., Tricco, A. C., Welch, V. A., Whiting, P., and Moher, D. (2021). The PRISMA 2020 statement: An updated guideline for reporting systematic reviews. *BMJ*, 372:n71.

Ralph, P. et al. (2020). Empirical standards for software engineering research.

Rao, D. and Callison-Burch, C. (2026). BibTeX citation errors in scientific publishing agents: Evaluation and mitigation. In *Conference on Language Modeling (COLM 2026)*.

Samsonau, S. V. (2026). sciwrite-lint: Verification infrastructure for the age of science vibe-writing.

Santesso, N., Glenton, C., Dahm, P., Garner, P., Akl, E. A., Alper, B., Brignardello-Petersen, R., Carrasco-Labra, A., De Beer, H., Hultcrantz, M., Kuijpers, T., Meerpohl, J., et al. (2020). GRADE guidelines 26: informative statements to communicate the findings of systematic reviews of interventions. *Journal of Clinical Epidemiology*, 119:126–135.

Struthers, C., Harwood, J., de Beyer, J. A., Dhiman, P., Logullo, P., and Schlüssel, M. (2021). GoodReports: Developing a website to help health researchers find and use reporting guidelines. *BMC Medical Research Methodology*, 21:217.

Toulmin, S. E. (2003). *The Uses of Argument*. Cambridge University Press, updated edition. Originally published 1958.

Turner, L., Shamseer, L., Altman, D. G., Weeks, L., Peters, J., Kober, T., Dias, S., Schulz, K. F., Plint, A. C., and Moher, D. (2012). Consolidated standards of reporting trials (CONSORT) and the completeness of reporting of randomised controlled trials (RCTs) published in medical journals. *Cochrane Database of Systematic Reviews*, 2012(11):MR000030.

VANRA Development Group (2026). VANRA (verification of assertions in narrative review articles): a reporting checklist for narrative review articles. Reporting guidelines under development for systematic reviews, registered 2 September 2026. Contact: D. J. Parker. Guideline under development. Accessed 2026-09-26.

Veen, J. (2026). Agent-ready papers: verification infrastructure for AI-augmented academic writing. GitHub repository, CC BY 4.0.

Wachsmuth, H., Naderi, N., Hou, Y., Bilu, Y., Prabhakaran, V., Alberdingk Thijm, T., Hirst, G., and Stein, B. (2017). Computational argumentation quality assessment in natural language. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 1, Long Papers*, pages 176–187. Association for Computational Linguistics.

Walters, W. H. and Wilder, E. I. (2023). Fabrication and errors in the bibliographic citations generated by ChatGPT. *Scientific Reports*, 13:14045.

Walton, D., Reed, C., and Macagno, F. (2008). *Argumentation Schemes*. Cambridge University Press.

Whetten, D. A. (1989). What constitutes a theoretical contribution? *Academy of Management Review*, 14(4):490–495.

Xiong, M., Hu, Z., Lu, X., Li, Y., Fu, J., He, J., and Hooi, B. (2024). Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs. In *The Twelfth International Conference on Learning Representations (ICLR)*.

Yang, R. (2022). An empirical study on the scaffolding Chinese university students' English argumentative writing based on Toulmin model. *Heliyon*, 8(12):e12199.

Zhou, K., Hwang, J., Ren, X., and Sap, M. (2024). Relying on the unreliable: The impact of language models' reluctance to express uncertainty. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3623–3643.

Zhou, Y. and Yu, C. (2026). Auditable AI-assisted research writing: An engineering discipline with pre-registered process observation.