

The Verification Gap: Why AI-Augmented Academic Writing Needs Reporting Guidelines for Reasoning

[Author Name]¹

¹[Affiliation]

[email]

Abstract

Large language models are increasingly used to assist with academic writing. While citation hallucination has received considerable attention, AI writing assistants create additional verification challenges—confidence inflation, scope creep, and fluent prose that obscures reasoning gaps—that existing tools do not address. This gap is compounded by a structural absence in reporting guidelines: the EQUATOR Network maintains nearly 700 guidelines for empirical health research, yet no equivalent exists for non-empirical paper types. Drawing on Gregor’s (2006) theory-type taxonomy, we identify two fundamentally non-empirical theory types for which argument quality, rather than factual accuracy, is the primary verification challenge. We propose a typed verification model distinguishing claims (source-verifiable facts), arguments (Toulmin-verified reasoning chains), and propositions (Whetten-verified recommendations with boundary conditions), each requiring distinct procedures. Preliminary evidence from three retrospective audits suggests that typed verification reveals actionable issues invisible to conventional approaches, including false failures from type misclassification. We propose that the scholarly community develop verification infrastructure for non-empirical papers analogous to what EQUATOR provides for empirical research.

Keywords: academic writing, artificial intelligence, reporting guidelines, verification, argumentation, non-empirical research

Key Points:

- AI writing tools create verification challenges beyond citation hallucination, including confidence inflation and argument degradation.
- EQUATOR provides nearly 700 guidelines for empirical research; non-empirical papers have no equivalent infrastructure.
- We propose typed verification distinguishing claims, arguments, and propositions, each with distinct procedures.
- Retrospective audits show that wrong-type verification produces false failures invisible to conventional approaches.

1 Introduction: The Problem

The rapid adoption of large language models (LLMs) as writing assistants is transforming how academic papers are produced. Liang et al. (6) found that up to 22% of computer science papers now show evidence of LLM modification, with detectable usage across all major disciplines. These tools can synthesise literature, structure arguments, generate prose, and format references with remarkable fluency. Yet this fluency introduces risks that are distinct from—and in some ways more subtle than—the challenges of traditional academic writing. Where a human author’s limitations are typically visible in the text (awkward phrasing, gaps in coverage, inconsistent formatting), an AI-assisted draft can be polished and comprehensive while harbouring verification failures that are invisible to casual inspection.

The most widely discussed risk is citation hallucination. LLMs can invent plausible-sounding papers, complete with fabricated authors, journal names, and digital object identifiers. Mugaanyi et al. (7) found that 62–89% of DOIs generated by large language models across natural sciences and humanities topics were fabricated, with DOI accuracy as low as 8.5% in some disciplines. Unlike factual errors that contradict known evidence, hallucinated citations create a distinctive category of falsehood: references that cannot be falsified through domain knowledge alone because they refer to works that do not exist. Citation hallucination has received considerable attention, and tools designed to detect it are emerging.

However, citation hallucination may be only the most visible failure mode. A second, potentially more insidious risk is what we term *confidence inflation*: AI writing assistants may state speculative claims with the same certainty as well-established findings, using language such as “demonstrates” or “confirms” where “suggests” or “may indicate” would be more appropriate. In academic writing, the distinction between “demonstrates” and “suggests” carries substantial epistemic weight—it signals to readers whether a claim rests on strong evidence or preliminary findings. When AI-generated prose conflates these registers, the resulting text can misrepresent the state of evidence without containing any factually incorrect statements. Liang et al. (5) found that LLM-generated feedback

on research papers catches surface issues but struggles with the kind of deep argument analysis that would detect such confidence miscalibration. A retrospective audit of a technology paper under review at an IEEE journal identified six of twenty-two verifiable entries where the manuscript language was more confident than the underlying evidence warranted.

A third risk is scope creep. Without structural constraints such as section specifications and word budgets, AI-assisted drafts may expand arguments beyond what the available evidence supports, introduce tangential sections, and exceed format requirements. While scope management is a challenge in all academic writing, AI tools can amplify it by generating additional content with minimal effort, making expansion the path of least resistance.

These failure modes share a common characteristic: they cannot be detected by checking whether citations are real. Citation verification tools that automatically check whether references exist, along with model-level approaches such as retrieval-augmented generation (in which the model is given access to source documents during text generation), address the factual accuracy dimension of AI-assisted writing. These are valuable contributions, but they operate at the model or tool level—they improve what the AI generates or verify whether specific references exist. To our knowledge, no process-level verification infrastructure addresses the writing workflow itself: the systematic registration, tracking, and confidence calibration of claims throughout the drafting process.

The distinction between tool-level and process-level verification matters because the failure modes described above are not primarily failures of the AI model. They are failures of the writing process—failures that occur when there is no infrastructure to register claims at appropriate confidence levels, calibrate language to evidence strength, or verify that arguments are logically sound rather than merely well-sourced. This paper asks what such process-level infrastructure might look like and whether existing scholarly frameworks provide a foundation for building it.

2 The Landscape Gap

To assess the current state of verification infrastructure for academic writing, it is useful to examine the most comprehensive effort to date: the EQUATOR Network. EQUATOR (Enhancing the QUALity and Transparency Of health Research) maintains a database of nearly 700 reporting guidelines for health research (EQUATOR Network). These guidelines—including CONSORT for randomised controlled trials, STROBE for observational studies, and PRISMA for systematic reviews (8)—provide structured checklists that authors complete during manuscript preparation. The GoodReports tool (10) automates guideline selection based on study type, further reducing the barrier to adoption. A Cochrane review found that journals endorsing CONSORT showed improved reporting completeness across the majority of checklist items (13), and incomplete reporting has been identified as a major contributor to research waste (2). These guidelines represent the most mature infrastructure for structured verification in academic publishing.

However, an examination of the EQUATOR database reveals a structural gap. Virtually all of its nearly 700 guidelines address empirical research: studies that collect, analyse, and report data. Systematic reviews, covered by PRISMA (8), represent a partial exception—they are structured and methodical but synthesise rather than generate primary data. Beyond this partial exception, to our knowledge, no EQUATOR or equivalent guidelines exist for theoretical papers, design science contributions, perspective articles, or methodological papers that do not report empirical findings.

This absence is not accidental. It reflects a fundamental difference in what needs to be verified across paper types. Gregor (3) identified five theory types in information systems research based on four goals: analysis, explanation, prediction, and prescription. Two of these types—Type I (analytic theories, including taxonomies and frameworks) and Type V (design and action theories, including design principles and methods)—are fundamentally non-empirical in their primary contribution, though Type V papers may include empirical evaluation of the designed artefact. Their core contributions are not data points that can be verified against measurements or observations; they are conceptual structures whose value lies in their clarity, coherence, and utility. The remaining

types (Type II: explanatory, Type III: predictive, Type IV: explanatory and predictive) align more naturally with empirical verification approaches, as they make claims about observable phenomena.

While Gregor’s taxonomy was developed for information systems, the underlying distinction applies broadly. The IS community has developed structured methodologies for non-empirical work—notably Peffers et al.’s (2007) widely adopted design science research process model—yet these methodological frameworks have not been translated into reporting guidelines analogous to EQUATOR. More broadly, perspective articles, commentaries, theoretical syntheses, and methodological proposals all share the characteristic that their primary contribution is an argument or a framework rather than empirical evidence. These paper types appear across virtually every academic discipline, from the humanities and social sciences to engineering and medicine.

For these non-empirical paper types, the central verification challenge may not be “do the data support the claims?” but rather “is the argument logically sound, are the premises established, and are the boundary conditions specified?” (12; 14). This is a fundamentally different verification task from the one that EQUATOR guidelines are designed to support.

This suggests that the EQUATOR gap is structural rather than a mere oversight. Empirical papers need verification of research design, data collection, analysis, and reporting—domains for which structured checklists are well-suited. Non-empirical papers need verification of argument quality: warrant validity, reasoning coherence, confidence calibration, and boundary specification. These are different verification tasks that may require different infrastructure. While this gap is most acute in purely non-empirical papers, it represents a systemic risk wherever authors move from data-driven claims to interpretive reasoning—including the discussion sections of empirical research. The question is whether established frameworks from argumentation theory and theory-building methodology can provide a foundation for this infrastructure.

3 Typed Verification: A Proposal

We propose that verification infrastructure for non-empirical papers should distinguish three unit types, each requiring a distinct verification procedure (Table 1).

Table 1: Three verification unit types with per-type procedures.

Type	Definition	Verification question	Framework
Claim	Source-verifiable fact	Does the source say this?	Citation checking
Argument	Evidence + inference $\rightarrow$ conclusion	Is the reasoning valid?	Toulmin (1958/2003)
Proposition	Recommendation with boundaries	Are conditions specified?	Whetten (1989)

The first and most familiar type is the **claim**: a factual statement that can be verified against one or more sources. Claims are the default unit in citation-based verification—“Study X found Y” is a claim whose accuracy can be checked by reading Study X. Most existing verification tools, including citation checkers and fact-verification systems, operate on claims. For claims, the verification question is straightforward: does the cited source exist, and does it say what the manuscript attributes to it?

The second type is the **argument**: a reasoning chain that combines evidence with inference to reach a conclusion. Arguments appear throughout academic writing, particularly in discussion and conclusion sections, yet they resist claim-level verification because their validity depends not on any single source but on the logical relationship between premises and conclusion. Toulmin (12) provides a well-established framework for analysing argument structure, decomposing any argument into six components: claim (the conclusion asserted), grounds (the evidence offered), warrant (the inferential bridge connecting grounds to claim), backing (support for the warrant), qualifier (the degree of certainty), and rebuttal (conditions under which the argument does not hold). This framework may provide a basis for operationalising argument verification in academic writing: an argument can be evaluated by checking whether its grounds are established, its warrant is explicit and valid for the target audience, and its qualifier is calibrated to the strength of the evidence. (Our operationalisation focuses on five of these six components, treating backing as implicit when the warrant itself is accepted by the target audience.) Recent work by Gupta et al. (4) found that LLMs can extract Toulmin-style

argument structures through zero-shot prompting, suggesting that automated support for argument-level verification may be technically feasible.

The third type is the **proposition**: a novel recommendation or contribution that specifies what should be done and under what conditions. Propositions are the defining unit of many non-empirical papers—they are what the paper contributes to the field. Whetten (14) provides a framework for evaluating theoretical contributions through four components: What (the constructs or factors involved), How (the relationships between them), Why (the underlying logic or reasoning), and Who/Where/When (the boundary conditions specifying where the contribution applies). Propositions can be evaluated by checking whether all key constructs are defined, whether the reasoning is explicit, and whether boundary conditions are stated—without requiring empirical data.

These three types may benefit from a shared confidence calibration system. We propose mapping confidence tiers to prescribed language (Table 2), so that the certainty expressed in the text is systematically aligned with the strength of the underlying evidence. This mapping may enable authors and reviewers to detect confidence inflation—the failure mode described in Section 1—at the drafting stage rather than during peer review.

Table 2: Proposed confidence-to-language mapping.

Tier	Assign when...	Prescribed language
Established	Multiple independent sources; no dispute	“demonstrates,” “confirms”
Supported	2–3 sources agree; open questions remain	“indicates,” “evidence suggests”
Emerging	1–2 sources; not yet replicated	“may,” “preliminary evidence”
Speculative	Logical inference; no direct data	“warrants investigation”

The central argument of this paper is that these three types require fundamentally different verification procedures, and that applying the wrong procedure to the wrong type may produce systematic errors. A claim evaluated as an argument (checking its warrant rather than its source) wastes effort. More consequentially, an argument evaluated as a claim (checking whether it has a source rather than whether its reasoning is valid) may produce false failures—entries scored as unverified not because they are wrong, but because the verification procedure was inappropriate for the unit type. This argument rests on the premise that different epistemic contributions (facts, inferences, recommen-

dations) have different validity criteria; if this premise holds, then a unified verification procedure will systematically misclassify some entries. One might question whether a three-type model is necessary—whether a simpler binary distinction between empirical and non-empirical verification would suffice. However, the distinction between arguments and propositions matters because they fail differently: an argument with a weak warrant needs better reasoning, while a proposition with missing boundary conditions needs better scoping. Conflating the two under a single “non-empirical” category obscures which corrective action is appropriate. Preliminary evidence from the retrospective audits described in the following section suggests that this type of misclassification may occur in practice.

4 Preliminary Evidence

To explore whether typed verification reveals issues that conventional approaches miss, we retrospectively applied the framework described in Section 3 to three paper projects that had been developed with AI writing assistance but without typed verification infrastructure. These audits should be understood as proof of concept rather than validation—they are retrospective, conducted by the framework’s developers, and drawn from a single domain cluster.

The most instructive case involved a proposition paper arguing that feedback accuracy, not manikin realism, determines CPR training effectiveness. This paper had been developed with a conventional claim registry that tracked twenty-one entries, achieving 76% overall coverage and 100% coverage of critical claims. A retrospective audit reclassified all entries using the three-type model, revealing that sixteen were indeed claims, but two were arguments and three were propositions that had been misclassified as claims.

The consequences of this misclassification were most apparent for one entry—the paper’s central interpretive claim that invalid sensor feedback undermines the primary benefit of technology-enhanced feedback in CPR training. Under claim-level verification, this entry was scored at the lowest confidence tier (speculative, 0.25 on the project’s numeric scale) because it lacked a direct citation—it was classified as “logical inference.” However,

this entry is not a claim; it is an argument. Its three premises—that sensors exhibit a 5.6% calibration shift when integrated into a manikin, that datasheet specifications deviate 120–130 times from empirical calibration, and that most manikins lack metrological characterisation—are each independently verified. When evaluated using the Toulmin checklist, the entry achieves four of five verification criteria: the claim is clearly stated, the grounds are verified, the warrant is logically valid, and the qualifier is appropriately calibrated. Only counter-argument engagement is partial. This represents what we term a *false failure*: an entry scored as the weakest in the registry not because its reasoning is unsound, but because the wrong verification procedure was applied. The distinction matters for the manuscript: the writing guide mapped the speculative score to heavily hedged language (“warrants investigation”), which is excessively cautious for a logically sound argument whose premises are verified. Typed verification would have assigned an emerging confidence tier, permitting language such as “preliminary evidence suggests”—a material improvement in argumentative force.

A second retrospective audit applied confidence tier analysis to a technology paper under review at an IEEE instrumentation journal. This audit extracted twenty-two verifiable entries from the manuscript text and assigned each a confidence tier based on the available evidence. The analysis identified six entries where the manuscript language was more confident than the evidence warranted—for instance, using “demonstrates” for findings from a single validation study with a limited sample size.

A third audit applied the typed verification framework to an engineering characterisation paper targeting a measurement science journal. This project used a tag-based system that conflated verification status (“checked” versus “unchecked”) with confidence assessment (“how strong is the evidence?”). Retyping its twenty-one entries identified two arguments and three propositions that had been treated as claims, and the discussion section contained at least six substantive arguments that were not tracked in the registry at all.

Across all three retrospective audits, the typed verification framework revealed actionable issues in each project. These issues included mistyped entries producing false failures,

over-confident language, missing boundary conditions on propositions, and unregistered arguments in discussion sections. While the specific issues varied by paper type, the common finding was that each audit identified problems that were not visible without the typed verification infrastructure.

These preliminary findings are consistent with independent research on structured verification for AI-assisted scholarly work. Sukpanichnant et al. (11) developed PeerArg, a framework that combines formal argumentation structure with LLM-based review analysis; their results suggest that structured representation of support and attack relations improves review aggregation compared to end-to-end LLM approaches. Similarly, Gupta et al. (4) found that prompting LLMs with Toulmin’s argument model significantly improves argument extraction compared to generic prompts. Both findings suggest that combining structured frameworks with LLM capabilities may outperform LLM capabilities alone—a principle consistent with the process-level verification approach proposed here.

Several important limitations apply to these preliminary findings. All three audits were retrospective—the framework was applied after the papers were drafted, not during writing. (Appendix A presents a complementary data point: the verification registry maintained during the writing of this paper, showing the framework applied prospectively.) All three projects share the same author team and fall within a related domain cluster (medical simulation and measurement science). The audits were conducted by the framework’s developers, introducing potential confirmation bias. These limitations mean that the evidence should be interpreted as preliminary indications rather than validation. Whether typed verification generalises to other authors, domains, and writing workflows remains an open question.

5 Call to Action

The preceding sections have traced a chain of reasoning. AI writing assistance creates failure modes beyond citation hallucination, including confidence inflation and argument-level degradation (Section 1). The most comprehensive verification infrastructure, EQUA-

TOR’s nearly 700 reporting guidelines, addresses empirical research exclusively; non-empirical paper types—which include two of Gregor’s five theory types—have no equivalent verification framework (Section 2). Established frameworks from argumentation theory (Toulmin) and theory-building methodology (Whetten) may provide a foundation for typed verification that distinguishes claims, arguments, and propositions (Section 3). Preliminary evidence from retrospective audits suggests that this approach can reveal actionable issues, including false failures caused by type misclassification (Section 4).

On the basis of this reasoning, we propose that the scholarly community develop verification infrastructure for non-empirical papers. Such infrastructure would extend the principle underlying EQUATOR—that structured reporting improves research quality—to paper types that currently have no systematic verification support. Specifically, we propose that verification procedures be differentiated by unit type: source verification for claims, Toulmin-based analysis for arguments, and Whetten-based analysis for propositions. Each type would have associated checklists, confidence calibration criteria, and language guidelines that authors can apply during manuscript preparation.

This proposal has specific boundary conditions. It applies to academic papers intended for peer-reviewed publication in fields where argument quality constitutes a primary contribution. It does not apply to informal writing, journalism, creative writing, or very short opinion pieces where the overhead of formal verification would exceed the benefit. Furthermore, the specific verification frameworks proposed here—Toulmin’s argument model and Whetten’s theoretical contribution framework—are grounded in Western academic argumentation traditions. Their applicability to scholarly traditions with different rhetorical conventions would need to be assessed rather than assumed.

One might object that journal peer review already provides adequate argument verification for non-empirical papers—that reviewers routinely evaluate logical coherence, boundary conditions, and counter-argument engagement. This observation has merit, and peer review remains essential. However, AI writing assistance may amplify the verification challenge beyond what traditional peer review was designed to handle. When AI tools can generate fluent, well-structured prose with minimal effort, the volume and

apparent quality of submitted manuscripts may increase while the ratio of well-reasoned to superficially coherent arguments decreases. Reviewers already face growing workloads; adding the burden of detecting AI-generated confidence inflation and implicit warrant failures may exceed what the current system can absorb. Process-level verification infrastructure would function as a complement to peer review, not a replacement—catching argument-level issues before submission so that reviewers can focus on substantive evaluation rather than structural verification.

Several directions warrant investigation. First, the typed verification model proposed here has been developed and applied within medical simulation and measurement science. Discipline-specific verification templates—adapting the three-type model to the conventions of different fields—would be needed to assess the framework’s generalisability. Second, integration with AI writing tools could enable real-time verification during the drafting process, using the structured checklists and confidence calibration criteria as prompting infrastructure for writing assistants. Third, empirical validation comparing the quality of papers written with and without typed verification infrastructure would provide the evidence base needed to move the proposal from preliminary to supported. Fourth, while this paper focuses on non-empirical paper types, the discussion sections of empirical papers involve the same argumentative move from data to interpretation—and may benefit from the same typed verification procedures applied to that transition.

The EQUATOR Network has spent two decades building nearly 700 reporting guidelines for empirical health research. These guidelines have measurably improved the quality and transparency of published research (13; 2). To our knowledge, essentially no equivalent guidelines exist for non-empirical papers—theoretical contributions, design science research, perspective articles, and methodological proposals. As AI writing tools make it easier to produce fluent academic prose, the gap between what is published and what is rigorously verified may widen. We suggest it is time to begin building the verification infrastructure for the rest.

Acknowledgements

[To be added.]

A Self-Application of the Typed Verification Framework

This paper uses the framework it describes. Unlike the three retrospective audits in Section 4, the claim registry for this manuscript was maintained during writing—entries were registered, typed, and verified as the paper was drafted. This appendix summarises that registry to provide a prospective data point that reviewers can independently verify.

Registry summary

Table 3 presents all 18 registered entries. The distribution across types—15 claims, 2 arguments, 1 proposition—reflects the paper’s structure: a perspective article grounded primarily in verifiable facts and reasoning chains, culminating in a single recommendation.

Table 3: Complete verification registry for this manuscript (18 entries).

ID	Type	Priority	Confidence	Status
S1-1	Claim	P0	Supported	Verified
S1-2	Claim	P0	Supported	Verified
S1-3	Claim	P1	Emerging	Verified
S1-4	Claim	P0	Emerging	Verified
S2-1	Claim	P0	Established	Verified
S2-2	Claim	P0	Established	Verified
S2-3	Claim	P1	Established	Verified
S2-4	Claim	P1	Supported	Verified
S3-1	Claim	P1	Supported	Verified
S3-2	Claim	P1	Supported	Verified
S3-3	Claim	P1	Emerging	Verified
S3-4	Argument	P0	Emerging	Verified
S4-1	Claim	P0	Emerging	Verified
S4-2	Claim	P1	Emerging	Verified
S4-3	Claim	P1	Supported	Verified
S4-4	Argument	P1	Emerging	Verified
S5-1	Proposition	P0	Emerging	Verified
S5-2	Claim	P2	Speculative	Verified

Worked examples by type

Claim (S2-1). “The EQUATOR Network maintains nearly 700 reporting guidelines.” This entry was initially drafted as “approximately 500 guidelines” based on a secondary source. Verification against the EQUATOR website (accessed 3 March 2026) returned 699 guidelines, prompting a correction to “nearly 700.” The original figure would have understated the landscape by approximately 30%, illustrating how even a straightforward factual claim benefits from structured source checking.

Argument (S3-4). “Claims, arguments, and propositions require fundamentally different verification procedures.” This is the paper’s central argument. Table 4 shows the Toulmin decomposition as applied during verification.

Table 4: Toulmin decomposition of the central argument (S3-4).

Component	Content
Claim	Different unit types require different verification procedures.
Grounds	Proposition audit entry H4 was scored speculative (0.25) under claim-level verification but earned emerging under Toulmin analysis (4/5 criteria met, premises independently verified).
Warrant	If applying the wrong verification procedure to a correctly reasoned entry produces a false failure, then type-specific procedures are necessary to avoid systematic misclassification.
Qualifier	Emerging—demonstrated in one retrospective audit; not independently validated.
Rebuttal	The evidence rests on a single worked example; a larger sample across domains would be needed to establish generalisability.

Proposition (S5-1). “The scholarly community should develop verification infrastructure for non-empirical papers.” This is the paper’s single proposition. Its Whetten decomposition: *What*—typed verification infrastructure distinguishing claims, arguments, and propositions. *How*—per-type checklists adapted from Toulmin (arguments) and Whetten (propositions), with confidence-to-language mapping. *Why*—EQUATOR provides nearly 700 guidelines for empirical research (S2-1) while non-empirical papers have none (S2-2),

and AI writing amplifies the verification challenge. *Boundary conditions*—applies to academic papers intended for peer-reviewed publication; excludes informal writing, journalism, and creative writing; the proposed checklists are grounded in Western argumentation traditions and may need adaptation for other scholarly conventions.

Corrections discovered during verification

The verification process caught four categories of issue before submission: (1) the EQUATOR guideline count was corrected from approximately 500 to nearly 700 (S2-1); (2) the citation hallucination claim (S1-1) initially lacked a source, prompting a literature search that identified Mugaanyi et al. (7); (3) the retyping count for the proposition audit was corrected from 15/2/4 to 16/2/3 to match the audit source table (S4-1); and (4) six confidence calibration issues were identified in the technology audit (S4-2), where manuscript language was more confident than the evidence warranted. These corrections—ranging from a factual error to missing evidence to miscalibrated language—are the kinds of issues the framework is designed to catch.

References

- [EQUATOR Network] EQUATOR Network. Enhancing the QUALity and Transparency of health research. <https://www.equator-network.org/>. Accessed: 2026-03-03.
- [2] Glasziou, P., Altman, D. G., Bossuyt, P., Boutron, I., Clarke, M., Julious, S., Michie, S., Moher, D., and Wager, E. (2014). Reducing waste from incomplete or unusable reports of biomedical research. *The Lancet*, 383(9913):267–276.
- [3] Gregor, S. (2006). The nature of theory in information systems. *MIS Quarterly*, 30(3):611–642.
- [4] Gupta, A., Zuckerman, E., and O’Connor, B. (2024). Harnessing Toulmin’s theory for zero-shot argument explication. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL 2024)*, pages 10259–10276.

- [5] Liang, W., Zhang, Y., Cao, H., Wang, B., Ding, D. Y., Yang, X., Vodrahalli, K., He, S., Smith, D. S., Yin, Y., McFarland, D. A., and Zou, J. (2024). Can large language models provide useful feedback on research papers? A large-scale empirical analysis. *NEJM AI*, 1(8).
- [6] Liang, W., Zhang, Y., Wu, Z., Lepp, H., Ji, W., Zhao, X., Cao, H., Liu, S., He, S., Huang, Z., Yang, D., Potts, C., Manning, C. D., and Zou, J. Y. (2025). Quantifying large language model usage in scientific papers. *Nature Human Behaviour*, 9(12):2599–2609.
- [7] Mugaanyi, J., Cai, L., Cheng, S., Lu, C., and Huang, J. (2024). Evaluation of large language model performance and reliability for citations and references in scholarly writing: Cross-disciplinary study. *Journal of Medical Internet Research*, 26:e52935.
- [8] Page, M. J., McKenzie, J. E., Bossuyt, P. M., Boutron, I., Hoffmann, T. C., Mulrow, C. D., Shamseer, L., Tetzlaff, J. M., Akl, E. A., Brennan, S. E., Chou, R., Glanville, J., Grimshaw, J. M., Hróbjartsson, A., Lalu, M. M., Li, T., Loder, E. W., Mayo-Wilson, E., McDonald, S., McGuinness, L. A., Stewart, L. A., Thomas, J., Tricco, A. C., Welch, V. A., Whiting, P., and Moher, D. (2021). The PRISMA 2020 statement: An updated guideline for reporting systematic reviews. *BMJ*, 372:n71.
- [9] Peffers, K., Tuunanen, T., Rothenberger, M. A., and Chatterjee, S. (2007). A design science research methodology for information systems research. *Journal of Management Information Systems*, 24(3):45–77.
- [10] Struthers, C., Harwood, J., de Beyer, J. A., Dhiman, P., Logullo, P., and Schlüssel, M. (2021). GoodReports: Developing a website to help health researchers find and use reporting guidelines. *BMC Medical Research Methodology*, 21:217.
- [11] Sukpanichnant, P., Rapberger, A., and Toni, F. (2024). PeerArg: Argumentative peer review with LLMs. In *Proceedings of NeLaMKRR@KR 2024*. arXiv:2409.16813.
- [12] Toulmin, S. E. (2003). *The Uses of Argument*. Cambridge University Press, updated edition. Originally published 1958.

- [13] Turner, L., Shamseer, L., Altman, D. G., Weeks, L., Peters, J., Kober, T., Dias, S., Schulz, K. F., Plint, A. C., and Moher, D. (2012). Consolidated standards of reporting trials (CONSORT) and the completeness of reporting of randomised controlled trials (RCTs) published in medical journals. *Cochrane Database of Systematic Reviews*, (11):MR000030.
- [14] Whetten, D. A. (1989). What constitutes a theoretical contribution? *Academy of Management Review*, 14(4):490–495.